

STAT646 - Homework 5

Daniel Osorio - dcosoriorh@tamu.edu

Department of Veterinary Integrative Biosciences

Texas A&M University

1. Consider a genome that consists of 1,000 bases. The genome is organized into 20 genes, each of length 50 bases. The first 50 bases in the genome correspond to gene one, the next 50 bases to gene two, etc. In a particular sample containing mRNA transcripts from this genome, suppose there are 25 mRNA copies of gene 5, 60 mRNA copies of gene 10, and 40 mRNA copies of gene 15; none of the other genes are expressed. We apply NGS to create sequencing reads of length 10 bases. We then align the reads back to the genome and use the resulting read counts to quantify expression for each gene. Follow the instructions below to conduct a simulation. For grading purposes, you need only provide your R code, the figure described in part (e), and the read count table described in part (f).

- (a) Create the simulated genome by:

```
> set.seed(101)
> genome <- paste(sample(c("A", "C", "T", "G"), 1000, replace = TRUE),
+ collapse = "")
```

- (b) Create the fragments to be sequenced. For transcript i , randomly pick a location in the transcript at which to break. Once a break has been created, you have two smaller fragments. If a fragment is less than our desired read length of 10, throw it away. If a fragment is of length 10, 11, . . . , 14, its first 10 bases will be sequenced (so, it does not need to be fragmented any further). If a fragment is of length 15 bases or greater, it will be fragmented again. Repeat this sequence until you are left only with short ($\text{length} \in [10, 14]$) fragments that are ready to be sequenced.
- (c) Create the sequenced reads. For fragment j , create a sequencing read by reading off the genomic sequence corresponding to the first 10 bases of this fragment, allowing for sequencing errors to occur in 5% of all base calls.
- (d) Align the reads to the genome. For this, use the `pairwiseAlignment` function from `Biostrings` to do local alignment. Reward a base match with a score of +1, penalize a base mismatch with a score of -1, and penalize a gap with a score of -10 (use `gapOpening = 0` and `gapExtension = -10` in `pairwiseAlignment`). Use the randomization based test to assign a p-value to the alignment of each read, and only keep the reads that align with p-values < 0.05 ; to keep the computation time down, just do $B = 100$ randomizations, instead of the $B = 1,000$ that we did in the classroom examples.
- (e) Tabulate the number of times each base in the genome was covered by one or more of the aligned reads. Create a dot plot to visualize the coverages. By ?dot plot, I just mean a traditional dot plot in statistics not the specialized plot we made for comparing two sequences. For example, suppose there are 14 reads that were aligned to overlap with base k in the genome. Then, you would put 14 dots, one above the

other, at the k th location along the horizontal axis. Figure 1 on page 4 shows an example, for a scenario in which there are just two genes.

- (f) Tabulate the number of reads that overlapped, at least partially, with each of the three expressed genes. For example, if 50 of the sequenced reads overlapped, at at least one base, with gene 5, the read count for gene 5 would be 50.
2. Context = “hammer” RNA-Seq data. These are read counts for ENSEMBL genes in rats under two conditions: (1) spinal nerve ligation (SNL), which induces chronic neuropathic pain, and (2) normal. We have four samples in each condition, two each at 2 weeks and 2 months after start (SNL was done at start in the SNL rats).

<http://www.ncbi.nlm.nih.gov/pubmed?term=20452967>

Use DESeq2 to analyze the data:

- (a) On the “rlog”-transformed counts:
 - i. Carry out hierarchical clustering, using Euclidean distance and complete linkage. Provide the resulting dendrogram and scree plots. How many clusters would you say there are? Interpret the clusters with respect to the meta data.
 - ii. Carry out principal component analysis to explore structure among the rats(columns). How much variation do the first two PCs account for? Interpret the eigenvectors with respect to the meta data.
- (b) As we have two factors (one each for treatment group (SNL vs. normal) and time (2 weeks vs. 2 months)), consider the following linear model for gene k ’s counts: